Date	Title	Authors	Link	Read order
	From the Deep Learning Papers Reading Roadmap		https://github.com/floodsung/Deep-Learning-Papers-Reading-Roadmap	
5/28/2015	Deep learning	Yann LeCun, Yoshua Bengio, Geoffrey Hinton	http://www.cs.toronto.edu/~hinton/absps/	1
6/9/2006	A fast learning algorithm for deep belief nets	Geoffrey Hinton, Simon Osindero, Yee-Whye Teh	http://www.cs.toronto.edu/~hinton/absps	2
6/28/2006	Reducing the dimensionality of data with neural networks	Geoffrey Hinton, Ruslan R. Salakhutdinov	http://www.cs.toronto.edu/~hinton/scie	3
12/3/2012	Imagenet classification with deep convolutional neural networks	Alex Krizhevsky, Ilya Sutskever, Geoffrey Hinton	http://papers.nips.cc/paper/4824-imag	4
9/4/2014	Very deep convolutional networks for large-scale image recognition	Karen Simonyan, Andrew Zisserman	https://arxiv.org/abs/1409.1556v6	5
12/16/2013	Network In Network	Min Lin, Qiang Chen, Shuicheng Yan	https://arxiv.org/abs/1312.4400	6
9/17/2014	Going deeper with convolutions	Christian Szegedy, Wei Liu, Yangqing Jia, Pierre Serma	https://arxiv.org/abs/1409.4842v1	7
12/10/2015	Deep residual learning for image recognition	Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun	https://arxiv.org/abs/1512.03385v1	8
10/15/2012	Deep neural networks for acoustic modeling in speech recognition: The shared view	Geoffrey Hinton, et al.	http://ieeexplore.ieee.org/xpls/abs_all.jsp	9
3/22/2013	Speech recognition with deep recurrent neural networks	Alex Graves, Abdel-rahman Mohamed, Geoffrey Hinton	https://arxiv.org/abs/1303.5778v1	10
6/21/2014	Towards End-To-End Speech Recognition with Recurrent Neural Networks	Alex Graves, Navdeep Jaitly	http://www.jmlr.org/proceedings/papers/	11
7/24/2015	Fast and accurate recurrent neural network acoustic models for speech recognition	Haşim Sak, et al.	https://arxiv.org/abs/1507.06947v1	12
12/8/2015	Deep speech 2: End-to-end speech recognition in english and mandarin	Dario Amodei, et al.	https://arxiv.org/abs/1512.02595v1	13
10/17/2016	Achieving Human Parity in Conversational Speech Recognition	W. Xiong, J. Droppo, X. Huang, F. Seide, M. Seltzer, A.	https://arxiv.org/abs/1610.05256v2	14
7/3/2012	Improving neural networks by preventing co-adaptation of feature detectors	Geoffrey E. Hinton, Nitish Srivastava, Alex Krizhevsky, Il	https://arxiv.org/abs/1207.0580v1	15
6/1/2014	Dropout: a simple way to prevent neural networks from overfitting	Nitish Srivastava, Geoffrey Hinton, Alex Krizhevsky, Ilya	http://www.cs.toronto.edu/~hinton/abs	16
2/11/2015	Batch normalization: Accelerating deep network training by reducing internal covaria	Sergey Ioffe, Christian Szegedy	https://arxiv.org/abs/1502.03167v3	17
7/21/2016	Layer normalization	Jimmy Lei Ba, Jamie Ryan Kiros, Geoffrey Hinton	https://arxiv.org/abs/1607.06450v1	18
2/9/2016	Binarized Neural Networks: Training Neural Networks with Weights and Activations	Matthieu Courbariaux, Itay Hubara, Daniel Soudry, Ran	https://arxiv.org/abs/1602.02830v3	19
8/18/2016	Decoupled neural interfaces using synthetic gradients	Max Jaderberg, Wojciech Marian Czarnecki, Simon Oser	https://arxiv.org/abs/1608.05343v2	20
11/19/2015	Net2Net: Accelerating learning via knowledge transfer	Tianqi Chen, Ian Goodfellow, Jonathon Shlens	https://arxiv.org/abs/1511.05641v4	21
3/5/2016	Network Morphism	Tao Wei, Changhu Wang, Yong Rui, Chang Wen Chen	https://arxiv.org/abs/1603.01670v3	22
6/17/2013	On the importance of initialization and momentum in deep learning	Ilya Sutskever, James Martens, George Dahl, Geoffrey I	http://www.jmlr.org/proceedings/papers/	23
12/22/2014	Adam: A method for stochastic optimization	Diederik Kingma, Jimmy Ba,	https://arxiv.org/abs/1412.6980v8	24
6/14/2016	Learning to learn by gradient descent by gradient descent	Marcin Andrychowicz, Misha Denil, Sergio Gomez, Matt	https://arxiv.org/abs/1606.04474v2	25
10/1/2015	Deep compression: Compressing deep neural network with pruning, trained quantiz	Song Han, Huizi Mao, William J. Dally	https://arxiv.org/abs/1510.00149v5	26
2/24/2016	SqueezeNet: AlexNet-level accuracy with 50x fewer parameters and< 1MB model s	Forrest N. Iandola, Song Han, Matthew W. Moskewicz, K	https://arxiv.org/abs/1602.07360v4	27
12/29/2011	Building high-level features using large scale unsupervised learning	Quoc V. Le, Marc'Aurelio Ranzato, Rajat Monga, Matthie	https://arxiv.org/abs/1112.6209v5	28
12/20/2013	Auto-encoding variational bayes	Diederik Kingma, Max Welling	https://arxiv.org/abs/1312.6114v10	29
6/10/2014	Generative adversarial nets	Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bo	https://arxiv.org/abs/1406.2661v1	30
11/19/2015	Unsupervised representation learning with deep convolutional generative adversaria	Alec Radford, Luke Metz, and Soumith Chintala	https://arxiv.org/abs/1511.06434v2	31
2/16/2015	DRAW: A recurrent neural network for image generation	Karol Gregor, Ivo Danihelka, Alex Graves, Danilo Jimeni	https://arxiv.org/abs/1502.04623v2	32
1/25/2016	Pixel recurrent neural networks	Aaron van den Oord, Nal Kalchbrenner, Koray Kavukcuo	https://arxiv.org/abs/1601.06759v3	33
6/16/2016	Conditional image generation with PixelCNN decoders	Aaron van den Oord, Nal Kalchbrenner, Oriol Vinyals, La	https://arxiv.org/abs/1606.05328v2	34
8/4/2013	Generating sequences with recurrent neural networks	Alex Graves	https://arxiv.org/abs/1308.0850v5	35
6/3/2014	Learning phrase representations using RNN encoder-decoder for statistical machin	Kyunghyun Cho, Bart van Merriënboer, Caglar Gulcehre	https://arxiv.org/abs/1406.1078v3	36
9/10/2014	Sequence to sequence learning with neural networks	Ilya Sutskever, Oriol Vinyals, Quoc V. Le	https://arxiv.org/abs/1409.3215v3	37
9/1/2014	Neural Machine Translation by Jointly Learning to Align and Translate	Dzmitry Bahdanau, Kyunghyun Cho, Yoshua Bengio	https://arxiv.org/abs/1409.0473v7	38
6/19/2015	A neural conversational model	Oriol Vinyals, Quoc Le	https://arxiv.org/abs/1506.05869v3	39
10/20/2014	Neural turing machines	Alex Graves, Greg Wayne, Ivo Danihelka	https://arxiv.org/abs/1410.5401v2	40
5/4/2015	Reinforcement Learning Neural Turing Machines	Wojciech Zaremba, Ilya Sutskever	http://arxiv.org/abs/1505.00521v3	41
10/13/2014	Memory networks	Jason Weston, Sumit Chopra, Antoine Bordes	http://arxiv.org/abs/1410.3916v11	42
3/31/2015	End-to-end memory networks	Sainbayar Sukhbaatar, Arthur Szlam, Jason Weston, Ro	http://arxiv.org/abs/1503.08895v5	43
6/9/2015	Pointer networks	Oriol Vinyals, Meire Fortunato, Navdeep Jaitly	http://arxiv.org/abs/1506.03134v2	44
10/27/2016	Hybrid computing using a neural network with dynamic external memory	Alex Graves, Greg Wayne, Malcolm Reynolds, Tim Hark	https://www.gwern.net/docs/ai/2016-gra	45
12/19/2013	Playing atari with deep reinforcement learning	Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Alex	https://arxiv.org/abs/1312.5602v1	46
2/26/2015	Human-level control through deep reinforcement learning	Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andr	https://www.google.com/url?sa=t&rct=	47
11/20/2015	Dueling network architectures for deep reinforcement learning	Ziyu Wang, Tom Schaul, Matteo Hessel, Hado van Hasse	http://arxiv.org/abs/1511.06581v3	48
2/4/2016	Asynchronous methods for deep reinforcement learning	Volodymyr Mnih, Adrià Puigdomènech Badia, Mehdi Mir	http://arxiv.org/abs/1602.01783v2	49
9/9/2015	Continuous control with deep reinforcement learning	Timothy P. Lillicrap, Jonathan J. Hunt, Alexander Pritzel	http://arxiv.org/abs/1509.02971v5	50
3/2/2016	Continuous Deep Q-Learning with Model-based Acceleration	Shixiang Gu, Timothy Lillicrap, Ilya Sutskever, Sergey L	http://arxiv.org/abs/1603.00748v1	51
2/19/2015	Trust region policy optimization	John Schulman, Sergey Levine, Philipp Moritz, Michael I	http://arxiv.org/abs/1502.05477v5	52
1/1/2016	Mastering the game of Go with deep neural networks and tree search	David Silver, Julian Schrittwieser, Karen Simonyan, Ioan	https://storage.googleapis.com/deepmi	53
7/2/2011	Deep Learning of Representations for Unsupervised and Transfer Learning	Yoshua Bengio	http://www.jmlr.org/proceedings/papers	54
3/15/2013	Lifelong Machine Learning Systems: Beyond Learning Algorithms	Daniel L. Silver, Qiang Yang, Lianghao Li	http://citeseerx.ist.psu.edu/viewdoc/de	55
3/9/2015	Distilling the knowledge in a neural network	Geoffrey Hinton, Oriol Vinyals, Jeff Dean	https://arxiv.org/abs/1503.02531v1	56
11/19/2015	Policy distillation	Andrei A. Rusu, Sergio Gomez Colmenarejo, Caglar Gul	http://arxiv.org/abs/1511.06295v2	57
11/29/2015	Actor-mimic: Deep multitask and transfer reinforcement learning	Emilio Parisotto, Jimmy Lei Ba, Ruslan Salakhutdinov	http://arxiv.org/abs/1511.06342v4	58
6/15/2016	Progressive neural networks	Andrei A. Rusu, Neil C. Rabinowitz, Guillaume Desjardin	http://arxiv.org/abs/1606.04671v3	59
12/11/2015	Human-level concept learning through probabilistic program induction	Brenden M. Lake, Ruslan Salakhutdinov, Joshua B. Ten	http://clm.utexas.edu/compjclub/wp-co	60
6/6/2015	Siamese Neural Networks for One-shot Image Recognition	Gregory Koch, Richard Zemel, Ruslan Salakhutdinov	http://www.cs.utoronto.ca/~gkoch/file	61
5/19/2016	One-shot Learning with Memory-Augmented Neural Networks	Adam Santoro, Sergey Bartunov, Matthew Botvinick, Da	http://arxiv.org/abs/1605.06065v1	62
6/13/2016	Matching Networks for One Shot Learning	Oriol Vinyals, Charles Blundell, Timothy Lillicrap, Koray	https://arxiv.org/abs/1606.04080v2	63
6/9/2016	Low-shot Visual Recognition by Shrinking and Hallucinating Features	Bharath Hariharan, Ross Girshick	http://arxiv.org/abs/1606.02819v4	64
4/21/2012	Joint Learning of Words and Meaning Representations for Open-Text Semantic Par	Antoine Bordes, Xavier Glorot, Jason Weston, Yoshua B	https://www.hds.utc.fr/~bordesan/doku	65
10/16/2013	Distributed representations of words and phrases and their compositionality	Tomas Mikolov, Ilya Sutskever, Kai Chen, Greg Corrado	http://arxiv.org/abs/1310.4546v1	66
6/24/2015	Ask Me Anything: Dynamic Memory Networks for Natural Language Processing	Ankit Kumar, Ozan Irsoy, Peter Ondruska, Mohit Iyyer, J	http://arxiv.org/abs/1506.07285v5	67
7/26/2015	Character-Aware Neural Language Models	Yoon Kim, Yacine Jernite, David Sontag, Alexander M. R	http://arxiv.org/abs/1508.06615v4	68
2/19/2015	Towards AI-Complete Question Answering: A Set of Prerequisite Toy Tasks	Jason Weston, Antoine Bordes, Sumit Chopra, Alexande	http://arxiv.org/abs/1502.05698v10	69
6/10/2015	Teaching Machines to Read and Comprehend	Karl Moritz Hermann, Tomáš Kočiský, Edward Grefenste	http://arxiv.org/abs/1506.03340v3	70
6/6/2016	Very Deep Convolutional Networks for Text Classification	Alexis Conneau, Holger Schwenk, Loïc Barrault, Yann L	http://arxiv.org/abs/1606.01781v2	71
6/6/2016	Bag of Tricks for Efficient Text Classification	Armand Joulin, Edouard Grave, Piotr Bojanowski, Tomas	http://arxiv.org/abs/1607.01759v3	72
12/5/2013	Deep neural networks for object detection	Christian Szegedy, Alexander Toshev, Dumitru Erhan	http://papers.nips.cc/paper/5207-deep	73
11/11/2013	Rich feature hierarchies for accurate object detection and semantic segmentation	Ross Girshick, Jeff Donahue, Trevor Darrell, Jitendra Ma	http://arxiv.org/abs/1311.2524v5	74
6/18/2014	Spatial pyramid pooling in deep convolutional networks for visual recognition	Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun	http://arxiv.org/abs/1406.4729v4	75
4/30/2015	Fast R-CNN	Ross Girshick	http://arxiv.org/abs/1504.08083v2	76
6/4/2015	Faster R-CNN: Towards real-time object detection with region proposal networks	Shaoqing Ren, Kaiming He, Ross Girshick, Jian Sun	http://arxiv.org/abs/1506.01497v3	77
6/8/2015	You only look once: Unified, real-time object detection	Joseph Redmon, Santosh Divvala, Ross Girshick, Ali Fa	http://arxiv.org/abs/1506.02640v5	78
12/8/2015	SSD: Single Shot MultiBox Detector	Wei Liu, Dragomir Anguelov, Dumitru Erhan, Christian S	http://arxiv.org/abs/1512.02325v5	79
5/20/2016	R-FCN: Object Detection via Region-based Fully Convolutional Networks	Jifeng Dai, Yi Li, Kaiming He, Jian Sun	http://arxiv.org/abs/1605.06409v2	80
11/14/2014	Fully convolutional networks for semantic segmentation	Jonathan Long, Evan Shelhamer, Trevor Darrell	http://arxiv.org/abs/1411.4038v2	81
12/22/2014	Semantic image segmentation with deep convolutional nets and fully connected crfs	Liang-Chieh Chen, George Papandreou, Iasonas Kokkir	http://arxiv.org/abs/1412.7062v4	82
6/20/2015	Learning to segment object candidates	Pedro O. Pinheiro, Ronan Collobert, Piotr Dollar	http://arxiv.org/abs/1506.06204v2	83
12/14/2015	Instance-aware semantic segmentation via multi-task network cascades	Jifeng Dai, Kaiming He, Jian Sun	http://arxiv.org/abs/1512.04412v1	84
3/29/2016	Instance-sensitive Fully Convolutional Networks	Jifeng Dai, Kaiming He, Yi Li, Shaoqing Ren, Jian Sun	https://arxiv.org/abs/1603.08678v1	85
3/20/2017	Mask R-CNN	Kaiming He, Georgia Gkioxari, Piotr Dollár, Ross Girshic	https://arxiv.org/abs/1703.06870v3	86
12/5/2013	Learning a deep compact image representation for visual tracking	Naiyan Wang, Dit-Yan Yeung	http://papers.nips.cc/paper/5192-learn	87
1/19/2015	Transferring rich feature hierarchies for robust visual tracking	Naiyan Wang, Siyi Li, Abhinav Gupta, Dit-Yan Yeung	http://arxiv.org/abs/1501.04587v2	88
12/1/2015	Visual tracking with fully convolutional networks	Lijun Wang, et al.	http://www.cv-foundation.org/openacc	89
4/6/2016	Learning to Track at 100 FPS with Deep Regression Networks	David Held, Sebastian Thrun, Silvio Savarese	http://arxiv.org/abs/1604.01802v2	90
6/30/2016	Fully-Convolutional Siamese Networks for Object Tracking	Luca Bertinetto, Jack Valmadre, João F. Henriques, And	http://arxiv.org/abs/1606.09549v2	91
7/12/2016	Beyond Correlation Filters: Learning Continuous Convolution Operators for Visual T	Martin Danelljan, Andreas Robinson, Fahad Shahbaz Kh	http://arxiv.org/abs/1608.03773v2	92
7/25/2016	Modeling and Propagating CNNs in a Tree Structure for Visual Tracking	Hyeonseob Nam, Mooyeol Baek, Bohyung Han	http://arxiv.org/abs/1608.07242v1	93

Date	Title	Authors	Link	Read order	
11/15/2010	Every picture tells a story: Generating sentences from images	Ali Farhadi, et al.	https://www.cs.cmu.edu/~afarhadi/pap	94	
6/20/2011	Baby talk: Understanding and generating image descriptions	Girish Kulkarni, et al.	http://tamaraberg.com/papers/generat	95	
11/17/2014	Show and tell: A neural image caption generator	Oriol Vinyals, Alexander Toshev, Samy Bengio, Dumitru	http://arxiv.org/abs/1411.4555v2	96	
11/17/2014	Long-term recurrent convolutional networks for visual recognition and description	Jeff Donahue, Lisa Anne Hendricks, Marcus Rohrbach, S	http://arxiv.org/abs/1411.4389v4	97	
12/7/2014	Deep visual-semantic alignments for generating image descriptions	Andrej Karpathy, Li Fei-Fei	http://arxiv.org/abs/1412.2306v2	98	
11/18/2014	From captions to visual concepts and back	Hao Fang, Saurabh Gupta, Forrest Iandola, Rupesh Sriv	http://arxiv.org/abs/1411.4952v3	99	
11/20/2014	Learning a recurrent visual representation for image caption generation	Xinlei Chen, C. Lawrence Zitnick	http://arxiv.org/abs/1411.5654v1	100	
12/20/2014	Deep captioning with multimodal recurrent neural networks (m-rnn)	Junhua Mao, Wei Xu, Yi Yang, Jiang Wang, Zhiheng Hu	http://arxiv.org/abs/1412.6632v5	101	
2/10/2015	Show, attend and tell: Neural image caption generation with visual attention	Kelvin Xu, Jimmy Ba, Ryan Kiros, Kyunghyun Cho, Aaro	http://arxiv.org/abs/1502.03044v3	102	
10/30/2014	Addressing the rare word problem in neural machine translation	Minh-Thang Luong, Ilya Sutskever, Quoc V. Le, Oriol Vir	http://arxiv.org/abs/1410.8206v4	103	
7/31/2015	Neural Machine Translation of Rare Words with Subword Units	Rico Sennrich, Barry Haddow, Alexandra Birch	http://arxiv.org/abs/1508.07909v5	104	
7/17/2015	Effective approaches to attention-based neural machine translation	Minh-Thang Luong, Hieu Pham, Christopher D. Manning	http://arxiv.org/abs/1508.04025v5	105	
3/19/2016	A Character-Level Decoder without Explicit Segmentation for Neural Machine Trans	Junyoung Chung, Kyunghyun Cho, Yoshua Bengio	http://arxiv.org/abs/1603.06147v4	106	
10/10/2016	Fully Character-Level Neural Machine Translation without Explicit Segmentation	Jason Lee, Kyunghyun Cho, Thomas Hofmann	http://arxiv.org/abs/1610.03017v3	107	
9/26/2016	Google's Neural Machine Translation System: Bridging the Gap between Human an	Yonghui Wu, Mike Schuster, Zhifeng Chen, Quoc V. Le,	http://arxiv.org/abs/1609.08144v2	108	
7/26/2015	A neural algorithm of artistic style	Leon A. Gatys, Alexander S. Ecker, Matthias Bethge	http://arxiv.org/abs/1508.06576v2	109	
9/12/2016	Generative Visual Manipulation on the Natural Image Manifold	Jun-Yan Zhu, Philipp Krähenbühl, Eli Shechtman, Alexei	http://arxiv.org/abs/1609.03552v3	110	x
3/5/2016	Semantic Style Transfer and Turning Two-Bit Doodles into Fine Artworks	Alex J. Champandard	http://arxiv.org/abs/1603.01768v1	111	
3/28/2016	Colorful Image Colorization	Richard Zhang, Phillip Isola, Alexei A. Efros	http://arxiv.org/abs/1603.08511v5	112	
3/27/2016	Perceptual losses for real-time style transfer and super-resolution	Justin Johnson, Alexandre Alahi, Li Fei-Fei	http://arxiv.org/abs/1603.08155v1	113	x
10/24/2016	A learned representation for artistic style	Vincent Dumoulin, Jonathon Shlens, Manjunath Kudlur	http://arxiv.org/abs/1610.07629v5	114	
11/23/2016	Controlling Perceptual Factors in Neural Style Transfer	Leon A. Gatys, Alexander S. Ecker, Matthias Bethge, Aa	http://arxiv.org/abs/1611.07865v2	115	
3/10/2016	Texture Networks: Feed-forward Synthesis of Textures and Stylized Images	Dmitry Ulyanov, Vadim Lebedev, Andrea Vedaldi, Victor	http://arxiv.org/abs/1603.03417v1	116	
	Literature review for guided audio generation project		https://github.com/khiner/guided-audio-generation/blob/master/docs/literature_review_notes.md		
12/22/2017	On Using Backpropagation for Speech Texture Generation and Voice Conversion	Jan Chorowski, Ron J. Weiss, Rif A. Saurous, Samy Ber	http://arxiv.org/abs/1712.08363v2	117	
6/29/2017	Audio spectrogram representations for processing with convolutional neural networ	L. Wyse	http://arxiv.org/abs/1706.09559v1	118	
10/31/2017	Audio style transfer	Eric Grinstein, Ngoc Duong, Alexey Ozerov, Patrick Pére	https://arxiv.org/abs/1710.11385v3	119	
11/29/2017	Time Domain Neural Audio Style Transfer	Parag K. Mital	https://arxiv.org/abs/1711.11160v1	120	
1/4/2018	Neural Style Transfer for Audio Spectograms	Prateek Verma, Julius O. Smith	https://arxiv.org/abs/1801.01589v1	121	
2/15/2018	"Style" Transfer for Musical Audio Using Multiple Time-Frequency Representations	Shaun Barry, Youngmoo Kim	https://openreview.net/pdf?id=BybQ7z	122	x
9/12/2016	WaveNet: A Generative Model for Raw Audio	Aaron van den Oord, Sander Dieleman, Heiga Zen, Kare	https://arxiv.org/abs/1609.03499v2	123	
9/14/2016	Convolutional Recurrent Neural Networks for Music Classification	Keunwoo Choi, George Fazekas, Mark Sandler, Kyungh	http://arxiv.org/abs/1609.04243v3	124	
9/29/2016	CNN Architectures for Large-Scale Audio Classification	Shawn Hershey, Sourish Chaudhuri, Daniel P. W. Ellis, .	http://arxiv.org/abs/1609.09430v2	125	
12/22/2016	SampleRNN: An Unconditional End-to-End Neural Audio Generation Model	Soroush Mehri, Kundan Kumar, Ishaan Gulrajani, Rithes	http://arxiv.org/abs/1612.07837v2	126	
3/29/2017	Tacotron: Towards End-to-End Speech Synthesis	Yuxuan Wang, RJ Skerry-Ryan, Daisy Stanton, Yonghui	http://arxiv.org/abs/1703.10135v2	127	
4/5/2017	Neural Audio Synthesis of Musical Notes with WaveNet Autoencoders	Jesse Engel, Cinjon Resnick, Adam Roberts, Sander Die	http://arxiv.org/abs/1704.01279v1	128	
11/1/2017	Uncovering Latent Style Factors for Expressive Speech Synthesis	Yuxuan Wang, RJ Skerry-Ryan, Ying Xiao, Daisy Stanto	http://arxiv.org/abs/1711.00520v1	129	
12/16/2017	Natural TTS Synthesis by Conditioning WaveNet on Mel Spectrogram Predictions	Jonathan Shen, Ruoming Pang, Ron J. Weiss, Mike Sch	https://arxiv.org/abs/1712.05884v2	130	
11/28/2017	Parallel WaveNet: Fast High-Fidelity Speech Synthesis	Aaron van den Oord, Yazhe Li, Igor Babuschkin, Karen S	http://arxiv.org/abs/1711.10433v1	131	
3/23/2018	Style Tokens: Unsupervised Style Modeling, Control and Transfer in End-to-End Sp	Yuxuan Wang, Daisy Stanton, Yu Zhang, RJ Skerry-Rya	https://arxiv.org/abs/1803.09017v1	132	
10/16/2018	Hierarchical Generative Modeling for Controllable Speech Synthesis	Wei-Ning Hsu, Yu Zhang, Ron J. Weiss, Heiga Zen, Yon	http://arxiv.org/abs/1810.07217v2	133	
6/10/2016	Improved Techniques for Training GANs	Tim Salimans, Ian Goodfellow, Wojciech Zaremba, Vicki	http://arxiv.org/abs/1606.03498v1	134	
7/12/2016	InfoGAN: Interpretable Representation Learning by Information Maximizing Generat	Xi Chen, Yan Duan, Rein Houthooft, John Schulman, Ilya	http://arxiv.org/abs/1606.03657v1	135	
9/11/2016	Energy-based Generative Adversarial Network	Junbo Zhao, Michael Mathieu, Yann LeCun	http://arxiv.org/abs/1609.03126v4	136	
11/15/2016	Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Net	Christian Ledig, Lucas Theis, Ferenc Huszar, Jose Caba	http://arxiv.org/abs/1609.04802v5	137	
11/21/2016	Image-to-Image Translation with Conditional Adversarial Networks	Phillip Isola, Jun-Yan Zhu, Tinghui Zhou, Alexei A. Efros	http://arxiv.org/abs/1611.07004v3	138	x
12/31/2016	NIPS 2016 Tutorial: Generative Adversarial Networks	Ian Goodfellow	http://arxiv.org/abs/1701.00160v4	139	
1/26/2017	Wasserstein GAN	Martin Arjovsky, Soumith Chintala, Léon Bottou	http://arxiv.org/abs/1701.07875v3	140	
3/31/2017	Improved Training of Wasserstein GANs	Ishaan Gulrajani, Faruk Ahmed, Martin Arjovsky, Vincen	http://arxiv.org/abs/1704.00028v3	141	
3/26/2017	Bayesian GAN	Yunus Saatchi, Andrew Gordon Wilson	http://arxiv.org/abs/1705.09558v3	142	
10/19/2017	Generative Adversarial Networks: An Overview	Antonia Creswell, Tom White, Vincent Dumoulin, Kai Aru	http://arxiv.org/abs/1710.07035v1	143	
3/30/2017	Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks	Jun-Yan Zhu, Taesung Park, Phillip Isola, Alexei A. Efros	http://arxiv.org/abs/1703.10593v6	144	
10/27/2017	Progressive Growing of GANs for Improved Quality, Stability, and Variation	Tero Karras, Timo Aila, Samuli Laine, Jaakko Lehtinen	http://arxiv.org/abs/1710.10196v3	145	
2/12/2018	Adversarial Audio Synthesis	Chris Donahue, Julian McAuley, Miller Puckette	http://arxiv.org/abs/1802.04208v3	146	x
2/23/2018	Efficient Neural Audio Synthesis	Nal Kalchbrenner, Erich Elsen, Karen Simonyan, Seb No	http://arxiv.org/abs/1802.08435v2	147	x
11/24/2017	StarGAN: Unified generative adversarial networks for multi-domain image-to-image	Yunjey Choi, Minje Choi, Munyoung Kim, Jung-Woo Ha,	http://arxiv.org/abs/1711.09020v3	149	x
3/21/2018	Self-Attention Generative Adversarial Networks	Han Zhang, Ian Goodfellow, Dimitris Metaxas, Augustus	https://arxiv.org/abs/1805.08318v1	150	x
6/26/2018	The challenge of realistic music generation: modelling raw audio at scale	Sander Dieleman, Aäron van den Oord, Karen Simonyar	https://arxiv.org/abs/1806.10474v1	151	x
6/2/2018	The relativistic discriminator: a key element missing from standard GAN	Alexia Jolicoeur-Martineau	https://arxiv.org/abs/1807.00734v3	152	
6/12/2018	The GAN Landscape: Losses, Architecture, Regularization, and Normalization	Karol Kurach, Mario Lucic, Xiaohua Zhai, Marcin Michals	https://arxiv.org/abs/1807.04720v2	153	
7/31/2018	Wasserstein GAN and Waveform Loss-based Acoustic Model Training for Multi-spe	Yi Zhao, Shinji Takaki, Hieu-Thi Luong, Junichi Yamagis	https://arxiv.org/abs/1807.11679v1	154	
7/24/2018	Voice Conversion with Conditional SampleRNN	Cong Zhou, Michael Horgan, Vivek Kumar, Cristina Vasc	https://arxiv.org/abs/1808.08311v1	155	
10/20/2017	Deep Voice 3: Scaling Text-to-Speech with Convolutional Sequence Learning	Wei Ping, Kainan Peng, Andrew Gibiansky, Sercan O. A	http://arxiv.org/abs/1710.07654v3	156	
7/19/2018	ClariNet: Parallel Wave Generation in End-to-End Text-to-Speech	Wei Ping, Kainan Peng, Jitong Chen	http://arxiv.org/abs/1807.07281v3	157	
7/20/2018	Fast spectrogram inversion using multi-head convolutional neural networks	Sercan O. Arik, Heewoo Jun, Gregory Diamos	http://arxiv.org/abs/1808.06719v2	158	
7/9/2018	Glow: Generative Flow with Invertible 1x1 Convolutions	Diederik P. Kingma, Prafulla Dhariwal	http://arxiv.org/abs/1807.03039v2	159	
10/31/2018	WaveGlow: A Flow-based Generative Network for Speech Synthesis	Ryan Prenger, Rafael Valle, Bryan Catanzaro	https://arxiv.org/abs/1811.00002v1	160	x
11/26/2018	GAN Dissection: Visualizing and Understanding Generative Adversarial Networks	David Bau, Jun-Yan Zhu, Hendrik Strobelt, Bolei Zhou, J	https://arxiv.org/abs/1811.10597v2	161	
12/12/2018	A Style-Based Generator Architecture for Generative Adversarial Networks	Tero Karras, Samuli Laine, Timo Aila	https://arxiv.org/abs/1812.04948v3	162	x
12/18/2018	Autoencoder Based Architecture For Fast & Real Time Audio Style Transfer	Dhruv Ramani, Samarjit Karmakar, Anirban Panda, Asac	https://arxiv.org/abs/1812.07159v2	163	
2/11/2019	Adversarial Generation of Time-Frequency Features with application in audio synthe	Andrés Marafioti, Nicki Holighaus, Nathanaël Perraudin,	https://arxiv.org/abs/1902.04072v1	164	
2/23/2019	GANSynth: Adversarial Neural Audio Synthesis	Jesse Engel, Kumar Krishna Agrawal, Shuo Chen, Ishaa	https://arxiv.org/abs/1902.08710v2	165	x
10/23/2018	SING: Symbol-to-Instrument Neural Generator	Alexandre Défossez, Neil Zeghidour, Nicolas Usunier, Le	https://arxiv.org/abs/1810.09785v1	166	
2/18/2019	SC-FEGAN: Face Editing Generative Adversarial Network with User's Sketch and C	Youngjoo Jo, Jongyoul Park	https://arxiv.org/abs/1902.06838v1	167	
9/28/2018	Large Scale GAN Training for High Fidelity Natural Image Synthesis	Andrew Brock, Jeff Donahue, Karen Simonyan	https://arxiv.org/abs/1809.11096v2	168	x
3/6/2019	High-Fidelity Image Generation With Fewer Labels	Mario Lucic, Michael Tschannen, Marvin Ritter, Xiaohua	https://arxiv.org/abs/1903.02271v1	169	
5/21/2019	Parallel Neural Text-to-Speech	Kainan Peng, Wei Ping, Zhao Song, Kexin Zhao	http://arxiv.org/abs/1905.08459v1	179	
6/4/2019	MelNet: A Generative Model for Audio in the Frequency Domain	Sean Vasquez, Mike Lewis	http://arxiv.org/abs/1906.01083v1	192	x
	More random DL papers				
2/5/2018	The Matrix Calculus You Need For Deep Learning	Terence Parr, Jeremy Howard	http://arxiv.org/abs/1802.01528v3	180	
12/2/2015	Rethinking the Inception Architecture for Computer Vision	Christian Szegedy, Vincent Vanhoucke, Sergey Ioffe, Jo	http://arxiv.org/abs/1512.00567v3	181	
2/23/2016	Inception-v4, Inception-ResNet and the Impact of Residual Connections on Learning	Christian Szegedy, Sergey Ioffe, Vincent Vanhoucke, Ale	http://arxiv.org/abs/1602.07261v2	182	
3/23/2016	Wide Residual Networks	Sergey Zagoruyko, Nikos Komodakis	http://arxiv.org/abs/1605.07146v4	183	
7/25/2016	Densely Connected Convolutional Networks	Gao Huang, Zhuang Liu, Laurens van der Maaten, Kilian	http://arxiv.org/abs/1608.06993v5	184	
11/16/2016	Aggregated Residual Transformations for Deep Neural Networks	Saining Xie, Ross Girshick, Piotr Dollár, Zhuowen Tu, Ka	http://arxiv.org/abs/1611.05431v2	186	
4/17/2017	MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications	Andrew G. Howard, Menglong Zhu, Bo Chen, Dmitry Kal	http://arxiv.org/abs/1704.04861v1	185	
7/4/2017	ShuffleNet: An Extremely Efficient Convolutional Neural Network for Mobile Devices	Xiangyu Zhang, Xinyu Zhou, Mengxiao Lin, Jian Sun	http://arxiv.org/abs/1707.01083v2	187	
1/13/2018	MobileNetV2: Inverted Residuals and Linear Bottlenecks	Mark Sandler, Andrew Howard, Menglong Zhu, Andrey Z	http://arxiv.org/abs/1801.04381v4	188	
7/30/2018	ShuffleNet V2: Practical Guidelines for Efficient CNN Architecture Design	Ningning Ma, Xiangyu Zhang, Hai-Tao Zheng, Jian Sun	http://arxiv.org/abs/1807.11164v1	189	
10/30/2018	DropBlock: A regularization method for convolutional networks	Golnaz Ghiasi, Tsung-Yi Lin, Quoc V. Le	http://arxiv.org/abs/1810.12890v1	190	
5/6/2019	Searching for MobileNetV3	Andrew Howard, Mark Sandler, Grace Chu, Liang-Chieh	http://arxiv.org/abs/1905.02244v2	191	
	Prior work in unpaired audio translation				
11/30/2017	Parallel-data-free voice conversion using cycle-consistent adversarial networks	Takuhiro Kaneko, Hirokazu Kameoka	http://arxiv.org/abs/1711.11293v2	170	
4/2/2018	High-quality nonparallel voice conversion based on cycle-consistent adversarial net	Fuming Fang, Junichi Yamagishi, Isao Echizen, Jaime L	http://arxiv.org/abs/1804.00425v1	171	

Date	Title	Authors	Link	Read order	
5/21/2018	A Universal Music Translation Network	Noam Mor, Lior Wolf, Adam Polyak, Yaniv Taigman	http://arxiv.org/abs/1805.07848v2	148	x
6/6/2018	StarGAN-VC: Non-parallel many-to-many voice conversion with star generative adv	Hirokazu Kameoka, Takuhiro Kaneko, Kou Tanaka, Nob	http://arxiv.org/abs/1806.02169v2	173	
7/6/2018	Singing Style Transfer Using Cycle-Consistent Boundary Equilibrium Generative Ad	Cheng-Wei Wu, Jen-Yu Liu, Yi-Hsuan Yang, Jyh-Shing I	http://arxiv.org/abs/1807.02254v1	174	
9/25/2018	WaveCycleGAN: Synthetic-to-natural speech waveform conversion using cycle-con	Kou Tanaka, Takuhiro Kaneko, Nobukatsu Hojo, Hirokaz	http://arxiv.org/abs/1809.10288v2	172	
9/29/2018	Modulated Variational auto-Encoders for many-to-many musical timbre transfer	Adrien Bitton, Philippe Esling, Axel Chemla-Romeu-Sant	https://arxiv.org/abs/1810.00222v1	175	x
121/22/2018	TimbreTron: A WaveNet(CycleGAN(CQT(Audio))) Pipeline for Musical Timbre Tran:	Sicong Huang, Qiyang Li, Cem Anil, Xuchan Bao, Sagee	http://arxiv.org/abs/1811.09620v2	176	x
4/5/2019	WaveCycleGAN2: Time-domain Neural Post-filter for Speech Waveform Generatior	Kou Tanaka, Hirokazu Kameoka, Takuhiro Kaneko, Nob	https://arxiv.org/abs/1904.02892v2	177	
3/1/2019	A Unified Neural Architecture for Instrumental Audio Tasks	Steven Spratley, Daniel Beck, Trevor Cohn	https://arxiv.org/abs/1903.00142v1	178	
	Random to-read list				
	Recent Advances in Convolutional Neural Networks: http://arxiv.org/abs/1512.07108 (12/22/2015)				
	Neural source-filter waveform models for statistical parametric speech synthesis				
	InverSynth: Deep Estimation of Synthesizer Parameter Configurations from Audio Signals				
	Transferring neural speech waveform synthesizers to musical instrument sound generation				
	Attention Is All You Need				
	Dynamic Routing Between Capsules				
	BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding				
	Universal Transformers				
	Neural Ordinary Differential Equations				
	Neural Processes				
	Conditional Neural Processes				
	Deep Reinforcement Learning: An Overview				
	A Brief Survey of Deep Reinforcement Learning				
	Model-Based Reinforcement Learning for Atari				
	NetAdapt: Platform-Aware Neural Network Adaptation for Mobile Applications				
	DARTS: Differentiable Architecture Search				
	Weight Normalization: A Simple Reparameterization to Accelerate Training of Deep Neural Networks: http://arxiv.org/abs/1602.07868 (2/25/2016)				
	Self-Normalizing Neural Networks: http://arxiv.org/abs/1706.02515 (6/8/2017)				
	Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour: http://arxiv.org/abs/1706.02677 (6/8/2017)				
	Don't Decay the Learning Rate, Increase the Batch Size: http://arxiv.org/abs/1711.00489 (11/1/2017)				
	Group Normalization: http://arxiv.org/abs/1803.08494 (3/22/2018)				
	A disciplined approach to neural network hyper-parameters: Part 1 -- learning rate, batch size, momentum, and weight decay: http://arxiv.org/abs/1803.09820 (4/24/2018)				
	Relational inductive biases, deep learning, and graph networks				
	A Comprehensive Survey on Graph Neural Networks				
	Graph Neural Networks: A Review of Methods and Applications				
	Relational recurrent neural networks				
	Unsupervised learning of disentangled and interpretable representations from sequential data				
	VideoFlow: A Flow-Based Generative Model for Video				
	Video-to-Video Synthesis				
	Understanding deep learning requires rethinking generalization: http://arxiv.org/abs/1611.03530 (11/10/2016)				
	Information Theory: A Tutorial Introduction: http://arxiv.org/abs/1802.05968 (2/16/2018)				
	On the Implicit Assumptions of GANs: http://arxiv.org/abs/1811.12402 (11/29/2018)				
	Deep Learning for Audio Signal Processing	Hendrik Purwins, Bo Li, Tuomas Virtanen, Jan Schlüter,	http://arxiv.org/abs/1905.00078v1		
	MIDI-to-audio (symbol-conditional audio generation)				
	Conditioning Deep Generative Raw Audio Models for Structured Automatic Music				
	Audio GPU				
	Real-Time Additive Synthesis with One Million Sinusoids Using a GPU				
	CUDA Accelerated Audio Digital Signal Processing for Real-Time Algorithms				
	REAL-TIME FORCE-BASED SOUND SYNTHESIS USING GPU PARALLEL COMPUTING				
	GPU-ACCELERATED MODAL PROCESSORS AND DIGITAL WAVEGUIDES				
	Audio Generation				
	WaveFlow				
	https://datascienceweekly.us3.list-manage.com/track/click?u=71a2b2a38789d4d25b738462f&id=960df0e938&e=9e42ffe2e8				
	OpenAI jukebox				